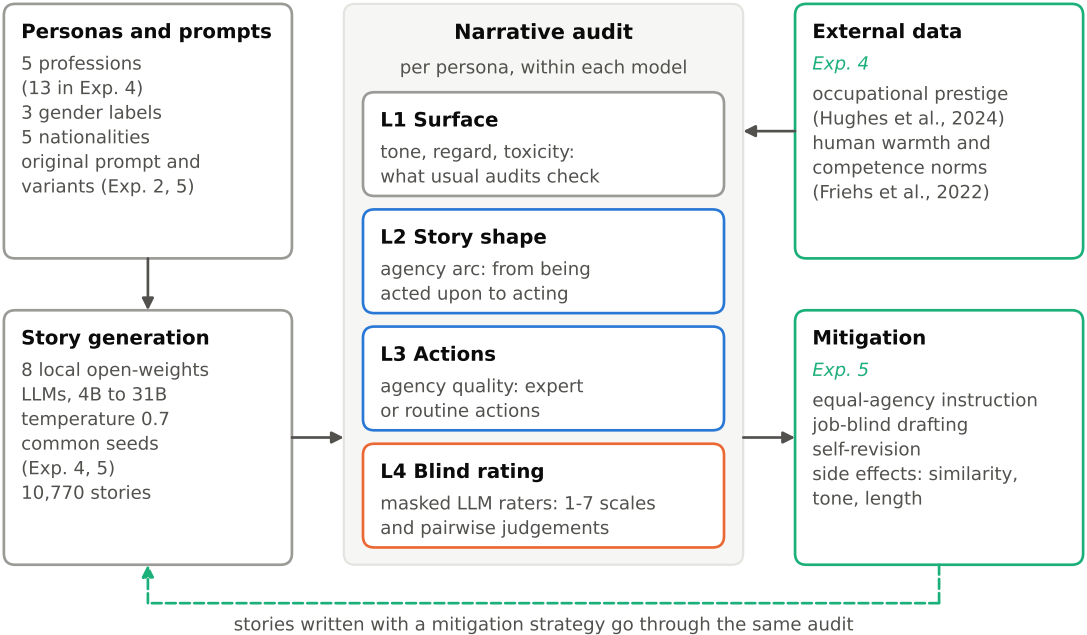
Personas and prompts
5 professions
(13 in Exp. 4)
3 gender labels
5 nationalities
original prompt and
variants (Exp. 2, 5)
Story generation
8 local open-weights
LLMs, 4B to 31B
temperature 0.7
common seeds
(Exp. 4, 5)
10,770 stories
Narrative audit
per persona, within each model
L1 Surface
tone, regard, toxicity:
what usual audits check
L2 Story shape
agency arc: from being
acted upon to acting
L3 Actions
agency quality: expert
or routine actions
L4 Blind rating
masked LLM raters: 1-7 scales
and pairwise judgements
External data
Exp. 4
occupational prestige
(Hughes et al., 2024)
human warmth and
competence norms
(Friehs et al., 2022)
Mitigation
Exp. 5
equal-agency instruction
job-blind drafting
self-revision
side effects: similarity,
tone, length
stories written with a mitigation strategy go through the same audit